


PROCESI I RASPOREDIVAČ

NA OS-U MOŽEMO IMATI VIŠE PROCESA (A I PARALELNO)

PROGRAM SA SVIM SVOJIM PODACIMA, FAJLOVI U MEMORIJI ČINI JEDAN PROCES



- XUG "MOŽE" DA RADI SA VIŠE PROCESORA
- TIKOV - JEDINICA VREMENA NA MATERNJOJ PLOČI

- DA BI SE OMOGUĆILO OVO "CIPERKANJE" PROCESA I NASTAVLJANJE MORAMO DA KORISTIMO SCHEDULER

RASPOREDIVAČ


XUG MOŽE DA IMA MAX 64 AKTIVNA PROCESA

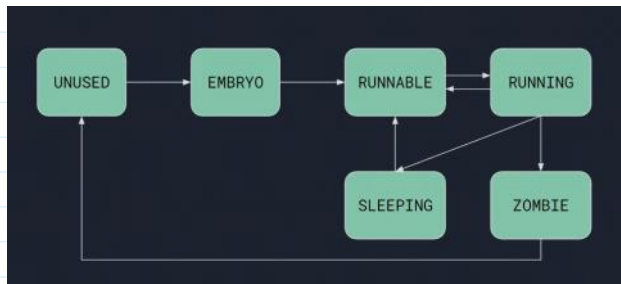
STRUKTURA PROC OPISUJE AKTIVNI PROCES

```
enum procstate { UNUSED, EMBRYO, SLEEPING, RUNNABLE, RUNNING, ZOMBIE };

// Per-process state
struct proc {
    uint sz; // Size of process memory (bytes)
    pde_t* pgdir; // Page table
    char *kstack; // Bottom of kernel stack for this process
    enum procstate state; // Process state
    int pid; // Process ID
    struct proc *parent; // Parent process
    struct trapframe *tf; // Trap frame for current syscall
    struct context *context; // switch() here to run process
    void *chan; // If non-zero, sleeping on chan
    int killed; // If non-zero, have been killed
    struct file *ofile[NOFILE]; // Open files
    struct inode *cwd; // Current directory
    char name[16]; // Process name (debugging)
};
```

CONTEXT SLUŽI ZA
POMOĆENJE NAPRETKA
JEDNOG PROCESA





ŽIVOTNI VIJEK PROCESA
U Xv6

UNUSED - PROCES KODI UOPŠTE NIJE PRI UPOTREBI
AKO ZA NEKI PROCES IMAMO PROCSTATE[UNUSED], TO
ZNAČI DA U NIZU OD 64 PROCESA ZA Xv6, MOŽEMO
DA GA ZAHVATIMO SA NEKIM DRUGIM PROCESOM
EMBRYO - NAGLAŠAVAMO DA ĆEMO PROCES DA KORISTIMO
RUNNABLE \longleftrightarrow RUNNING, PRELAZIMO IZ JEDNOG U DRUGI

- SCHEDULER KORISTI ROUND ROBIN ALGORITAM KROZ NIZ
PROCESA


- exit() FUNKCIJA STAVLJA RUNNING PROCES U ZOMBIE STANJE

- SVI PROCESI U Xv6 IMAJU PARENT-CHILD ODNOS.

- U TOJ HIERARHIJI, PRVI PROCES JESTE ZAPRAVO

INIT PROCES, POSLIJE/ISPOD NEGA IDE SHELL PROCES

\$



TRAP METODA JE HANDLER SVIM PREKIDA
TRAP.C

```

// Force process exit if it has been killed and is in user space.
// (If it is still executing in the kernel, let it keep running
// until it gets to the regular system call return.)
if(myproc() && myproc()->killed && (tf->cs&3) == DPL_USER)
    exit();

// Force process to give up CPU on clock tick.
// If interrupts were on while locks held, would need to check nlock.
if(myproc() && myproc()->state == RUNNING &&
    tf->strampno == T_IRQ0+IRQ_TIMER)
    yield();

// Check if the process has been killed since we yielded
if(myproc() && myproc()->killed && (tf->cs&3) == DPL_USER)
    exit();

```

```

struct {
    struct spinlock lock;
    struct proc proc[NPROC];
} ptable;

```

69 STAVKI

```

// Give up the CPU for one scheduling round.
void
yield(void)
{
    acquire(&ptable.lock); //DOC: yieldlock
    myproc()->state = RUNNABLE;
    sched();
    release(&ptable.lock);
}

```

PROC. C

```

void
sched(void)
{
    int intena;
    struct proc *p = myproc();

    if(!holding(&ptable.lock))
        panic("sched ptable.lock");
    if(mycpu()->ncli != 1)
        panic("sched locks");
    if(p->state != RUNNING)
        panic("sched running");
    if(readeflags() & FL_IF)
        panic("sched interruptible");
    intena = mycpu()->intena;
    switch(&p->context, mycpu()->scheduler);
    mycpu()->intena = intena;
}

```

USLOVI
KOD
SE NE TREBA DU
DEŠITI

SWITCH (ASSEMBLI)

```

# Context switch
#
# void switch(struct context **old, struct context *new);
#
# Save the current registers on the stack, creating
# a struct context, and save its address in *old.
# Switch stacks to new and pop previously-saved registers.

.globl switch
switch:
    movl 4(%esp), %eax
    movl 8(%esp), %edx

    # Save old callee-saved registers
    pushl %ebx
    pushl %esi
    pushl %edi

    # Switch stacks
    movl %esp, (%eax)
    movl %edx, %esp

    # Load new callee-saved registers
    popl %edi
    popl %esi
    popl %ebx
    popl %ebp
    ret

```

[ESP+4]
← OLD
← NEW

MITENJA NJE
UR DEANOST

POPUJENOM SA
SCHEDULERA

SRŽ
XUG
OPERATIVNOG
SISTEMA
JE BESKONAČNA
PETLJA

```

void
scheduler(void)
{
    struct proc *p;
    struct cpu *c = mycpu();
    c->proc = 0;

    for(;;){
        // Enable interrupts on this processor.
        sti();

        // Loop over process table looking for process to run.
        acquire(&ptable.lock);
        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
            if(p->state != RUNNABLE)
                continue;

            // Switch to chosen process. It is the process's job
            // to release ptable.lock and then reacquire it
            // before jumping back to us.
            c->proc = p;
            switchvm(p);
            p->state = RUNNING;

            switch(&c->scheduler, p->context);
            switchvm();

            // Process is done running for now.
            // It should have changed its p->state before coming back.
            c->proc = 0;

            release(&ptable.lock);
        }
    }
}

```

```

allocproc(void)
{
    struct proc *p;
    char *sp;

    acquire(&ptable.lock);

    for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
        if(p->state == UNUSED)
            goto found;
    }

    release(&ptable.lock);
    return 0;

found:
    p->state = EMULATED;
    p->pid = nextpid++;

    release(&ptable.lock);

    // Allocate kernel stack.
    if((p->kstack = kalloc()) == 0){
        p->state = UNUSED;
        return 0;
    }
    sp = p->kstack + KSTACKSIZE;

    // Leave room for trap frame.
    sp -= sizeof *p->tf;
    p->tf = (struct trapframe*)sp;

    // Set up new context to start executing at forkret,
    // which returns to trapret.
}

```

"POD ZNAČIČA NALAZA DA SE TELEPORTUJE"

PAUZA -11-

1:26:00

14-2008: Velje 10 5.4.2023. (raf.edu.rs)

```

Initc
int
main(void)
{
    int pid, wpid;


    if(getpid() != 1){
        printf(2, "init: already running\n");
        exit();
    }


    if(open("/dev/console", 0_RDONLY) < 0){
        mknod("/dev/console", 1, 1);
        open("/dev/console", 0_RDONLY);
    }

    dup(0); // stdout
    dup(0); // stderr

    for(;;){
        printf("init: starting sh\n");
        pid = fork();
        if(pid < 0){

```





```

dup(0); // stdout
dup(0); // stderr

for(;;){
    printf("init: starting sh\n");
    pid = fork();
    if(pid < 0){
        printf("init: fork failed\n");
        exit();
    }
    if(pid == 0){
        exec("/bin/sh", argv);
        printf("init: exec sh failed\n");
        exit();
    }
    while((wpid=wait()) >= 0 && wpid != pid)
        printf("zombie!\n");
}

```

Exit funkcija mnogo bitna

```

begin_op();
iput(curproc->cwd);
end_op();
curproc->cwd = 0;

acquire(&ptable.lock);

// Parent might be sleeping in wait().
wakeup1(curproc->parent);

// Pass abandoned children to init.
for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
    if(p->parent == curproc){
        p->parent = initproc;
        if(p->state == ZOMBIE)
            wakeup1(initproc);
    }
}

// Jump into the scheduler, never to return.
curproc->state = ZOMBIE;
sched();
panic("zombie exit");

```