

RADE SE SISTEMSKI POZIVI.

USER
!
KERNEL
IMAJU
STAKOV
METORIDE

USER SPACE

(UNUTAR NJEGA SE POZIVAJU
I IZVRŠAVAJU PROGRAMI IZ
USER FOLDERA)



SISTEMSKI POZIV KREIRA
SOFTVERSKI PREKID (INTERRUPT)

OVO JE ZID (KROZ NJEGA SE PROIBIJA POMOĆU


SISTEMSKIM POZIVIMA, ONI OMOGUĆAVAJU USER PROGRAMIMA DA KORISTE KERNEL FUNKCIJE



KERNEL IMA HANDLER PREKIDA
KOJI IH PREPOZNAJE

KERNEL SPACE

READ! WRITE!



XV6 IMA 2 REŽIMA RADA NORMALNI I ZAŠTILENI

FUNKCIJE UNUTAR KERNEL PROSTORA MOGU PARADATI KADA
GOĐU

USER.H
USYS.S
SYSCALL.H

```
1 #include "kernel/syscall.h"
2 #include "kernel/traps.h"
3
4 #define SYSCALL(name) \
5     .globl name; \
6     name: \
7     movl $SYS_ ## name, %eax; \
8     int $T_SYSCALL; \
9     ret;
10
11 SYSCALL(fork)
12 SYSCALL(exit)
13 SYSCALL(wait)
14 SYSCALL(pipe)
15 SYSCALL(read)
16 SYSCALL(write)
17 SYSCALL(close)
18 SYSCALL(kill)
19 SYSCALL(exec)
20 SYSCALL(open)
21 SYSCALL(mmap)
22 SYSCALL(unlink)
23 SYSCALL(fstat)
24 SYSCALL(link)
25 SYSCALL(mkdir)
26 SYSCALL(chdir)
27 SYSCALL(dup)
28 SYSCALL(getpid)
29 SYSCALL(sbrk)
30 SYSCALL(sleep)
31 SYSCALL(uptime)
```

ASSEMBLER U AT&T SINTAKSI

global open:
movl \$SYS_OPEN, %eax
int \$T_SYSCALL
ret
} MAKRO
SISTEMSKOG
POZIVA

SISTEMSKI PREKID IMA KOD 64

SISTEMSKI POZIV

KODISNIČKI PROGRAM POZIVUJE + USYS.S
NEKU METODU (OPEN()) NAJČEŠĆE IMAJU
ZA METODU I STAVLJA NA STEK

INT T_SYSCALL
VEŠTAČ. S → TRAP AS.M. S

DOLAZIMO DO TRAP.C
I TRAPFRAME STRUCTURE

syscall.c

sysfile.c

```
1 // System call numbers
2 #define SYS_fork 1
3 #define SYS_exit 2
4 #define SYS_wait 3
5 #define SYS_pipe 4
6 #define SYS_read 5
7 #define SYS_write 6
8 #define SYS_kill 7
9 #define SYS_exec 8
10 #define SYS_fstat 9
11 #define SYS_chdir 10
12 #define SYS_dup 11
13 #define SYS_getpid 12
14 #define SYS_sbrk 13
15 #define SYS_sleep 14
16 #define SYS_uptime 15
17 #define SYS_open 16
18 #define SYS_close 17
19 #define SYS_mmap 18
20 #define SYS_unlink 19
21 #define SYS_link 20
22 #define SYS_mkdir 21
```

```
void   
trap(struct trapframe *tf)   
{   
    if (tf->trapno == T_SYSCALL) {   
        if (myproc()->uilled)   
            exit();   
        myproc()->tf = tf;   
        syscall();   
        if (myproc()->uilled)   
            exit();   
    }   
}
```

```

    if (tf->trapno == T_SYSCALL) {
        if (myproc()-killed)
            exit();
        myproc()-tf = tf;
        syscall();
        if (myproc()-killed)
            exit();
        return;
    }

    switch (tf->trapno) {
        case T_IRQ0 + IRQ_TIMER:
            if (cpuid()) == 0 {
                acquire(&tickslock);
                ticks++;
                wakeup(&ticks);
                release(&tickslock);
            }
            lapicid();
            break;
    }
}

```

syscall.c

↓
sysfile.c

NUM CHECK →

```

void
syscall(void)
{
    int num;
    struct proc *curproc = myproc();

    num = curproc->tf->eax;
    if (num > 0 && num < N_SYSCALLS && syscall[num]) {
        curproc->tf->eax = syscall[num];
    } else {
        printf("%d %s: unknown sys call %d\n",
            curproc->pid, curproc->name, num);
        curproc->tf->eax = -1;
    }
}

```

ZADAN I AKTIVNI PROCES

```

// Per-process state
struct proc {
    void *sz; // Size of process memory (bytes)
    void *pgtbl; // Page table
    char *stack; // Bottom of kernel stack for this process
    enum procstate state; // Process state
    int pid; // Process ID
    struct proc *parent; // Parent process
    struct trapframe *tf; // Trap frame for current syscall
    struct context *context; // switch() here to run process
    void *kmem; // If non-zero, sleeping on kmem
    int killed; // If non-zero, have been killed
    struct file *ofile[NFILE]; // Open files
    struct inode *cwd; // Current directory
    char name[16]; // Process name (debugging)
};

```

SYSCALLS NIZ

```

static int (*syscalls[])(void) = {
[SYS_fork] sys_fork,
[SYS_exit] sys_exit,
[SYS_wait] sys_wait,
[SYS_pipe] sys_pipe,
[SYS_read] sys_read,
[SYS_kill] sys_kill,
[SYS_exec] sys_exec,
[SYS_fstat] sys_fstat,
[SYS_chdir] sys_chdir,
[SYS_dup] sys_dup,
[SYS_getpid] sys_getpid,
[SYS_sbrk] sys_sbrk,
[SYS_sleep] sys_sleep,
[SYS_uptime] sys_uptime,
[SYS_open] sys_open, ← 15.
[SYS_write] sys_write,
[SYS_mknod] sys_mknod,
[SYS_unlink] sys_unlink,
[SYS_link] sys_link,
[SYS_mkdir] sys_mkdir,
[SYS_close] sys_close,
};

```

UNUTAR SYSFILE.C SE NAHAZE FUNKCIJE POPUT SYS_OPEN(void)


```

204 int
205 sys_open(void)
206 {
207     char *path;
208     int fd, omode;
209     struct file *f;
210     struct inode *ip;
211
212     if (argstr(0, &path) < 0 || argint(1, &omode) < 0)
213         return -1;
214     begin_op();
215
216     if (omode & O_CREAT) {
217         ip = create(path, T_FILE, 0, 0);
218         if (ip == 0) {
219             end_op();
220             return -1;
221         }
222     } else {
223         if (ip = namei(path) == 0) {
224             end_op();
225             return -1;
226         }
227         if (ip->type == T_DIR && omode & O_RDONLY) {
228             iunlockput(ip);
229             end_op();
230             return -1;
231         }
232     }
233 }

```

→
OVO JE NEKA ARGUMENTI
ZA FUNKCIJU, ALI
PODACI KOJI NAM
TREBAJU SE NAHAZE
NA STEKU, Njih
DOBIVAMO KORISTEĆI
ARGINT I ARGSTR

KAKO TO FUNKCIONISE




```
// System call numbers
#define SYS_fork 1
#define SYS_exit 2
#define SYS_wait 3
#define SYS_pipe 4
#define SYS_read 5
#define SYS_kill 6
#define SYS_exec 7
#define SYS_fstat 8
#define SYS_chdir 9
#define SYS_dup 10
#define SYS_getpid 11
#define SYS_sbrk 12
#define SYS_sleep 13
#define SYS_uptime 14
#define SYS_open 15
#define SYS_write 16
#define SYS_mknod 17
#define SYS_unlink 18
#define SYS_link 19
#define SYS_mkdir 20
#define SYS_close 21
#define SYS_hello 22
```

ZAJIM KROZ
TRAP.C SE POZIVA
SISTEMSK, POZIV
ODGOVARAJUĆE
BROJ

OVDE POZIVU ODGOVARAJUĆE
BROJ

```
extern int sys_uptime(void);
extern int sys_hello(void);

static int (*syscalls[])(void) = {
[SYS_fork] sys_fork,
[SYS_exit] sys_exit,
[SYS_wait] sys_wait,
[SYS_pipe] sys_pipe,
[SYS_read] sys_read,
[SYS_kill] sys_kill,
[SYS_exec] sys_exec,
[SYS_fstat] sys_fstat,
[SYS_chdir] sys_chdir,
[SYS_dup] sys_dup,
[SYS_getpid] sys_getpid,
[SYS_sbrk] sys_sbrk,
[SYS_sleep] sys_sleep,
[SYS_uptime] sys_uptime,
[SYS_open] sys_open,
[SYS_write] sys_write,
[SYS_mknod] sys_mknod,
[SYS_unlink] sys_unlink,
[SYS_link] sys_link,
[SYS_mkdir] sys_mkdir,
[SYS_close] sys_close,
[SYS_hello] sys_hello,
}
```

SYS FILE.C

```
int sys_hello(void)
{
    int n;

    if(argInt(0, &n) < 0) return -1;

    int i;
    for(i = 0; i < n; i++)
    {
        cprintf("Hello world\n");
    }

    return 0;
}
```

UNUTAR

+

MAKEFILE
SCRIPT

ODRAZGO

KORISNIČKI PROGRAM

TO JE TO

FILE STRUKTURA

FILE.H

```
1 struct file {
2     enum { FD_NONE, FD_PIPE, FD_INODE } type;
3     int ref; // reference count
4     char readable;
5     char writable;
6     struct pipe *pipe;
7     struct inode *ip;
8     uint off;
9 };
10
```

RAD SA FAJLOVIMA U FILE.C I SYS FILE.C


FCNTL.H

```
1 #define O_RDONLY 0x000
2 #define O_WRONLY 0x001
3 #define O_RDWR 0x002
4 #define O_CREATE 0x200
5
```

IMAMO FUNKCIJE OPEN, READ, WRITE

FTABLE

FTABLE



```

sys_read(void)
{
    struct file *f;
    int n;
    char *p;

    if(argfd(0, 0, &f) < 0 || argint(2, &n) < 0 || argptr(1, &p, n) < 0)
        return -1;
    return fileread(f, p, n);
}
    
```

```

int
fileread(struct file *f, char *addr, int n)
{
    int r;

    if(f->readable == 0)
        return -1;
    if(f->type == FD_PIPE)
        return piperead(f->pipe, addr, n);
    if(f->type == FD_INODE){
        ilock(f->ip);
        if((r = readi(f->ip, addr, f->off, n)) > 0)
            f->off += r;
        iunlock(f->ip);
        return r;
    }
    panic("fileread");
}
    
```

```


int
readi(struct inode *ip, char *dst, uint off, uint n)
{
    uint tot, m;
    struct buf *bp;

    if(ip->type == T_DEV){
        if(ip->major < 0 || ip->major >= NDEV || !devsw[ip->major].read)
            return -1;
        return devsw[ip->major].read(ip, dst, n);
    }

    if(off > ip->size || off + n < off)
        return -1;
    if(off + n > ip->size)
        n = ip->size - off;

    for(tot=0; tot<n; tot+=m, off+=m, dst+=m){
        bp = bread(ip->dev, bmap(ip, off/BSIZE));
        m = min(n - tot, BSIZE - off%BSIZE);
        memmove(dst, bp->data + off%BSIZE, m);
        brelse(bp);
    }
    return tot;
}
    
```

OFF = OFFSET PROVIDED, MA
KOD SE WILREMENTIDG
ZA EITAN, E



BUFER

TOT → TOTAL NUMBER
OF BYTES READ FROM FILE/INODE

